

Evaluation report

Generated 2026-09-21T22:52:07 against the live n8n workflow. Model `openai/gpt-oss-120b`, prompts `{"classification": "classification.v2", "enrichment": "enrichment.v2", "summary": "summary.v1"}`.

21/23 checks passed. Latency per message (3 LLM calls): median 1987 ms, max 4080 ms.

Expectations are in `tests/golden.json` (written before running). A category list means more than one answer is acceptable.

Results

Case	Run	Expected category	Actual	Priority	Conf.	Expected queue	Destination	Expect esc.	Escalated (rules)	Result
S1	1	Bug Report	Bug Report	High	0.94	engineering	engineering	no	no (-)	PASS
S2	1	Feature Request	Feature Request	Low	0.94	product	product	no	no (-)	PASS
S3	1	Billing Issue	Billing Issue	Medium	0.95	billing	human_review_escalation	yes	yes (billing_amount_over_threshold)	PASS
S4	1	Technical Question	Technical Question	Low	0.94	technical_support	technical_support	either	no (-)	PASS
S5	1	Incident/Outage	Incident/Outage	High	0.96	incident_response	human_review_escalation	yes	yes (incident_category, outage_language)	PASS
S1	2	Bug Report	Bug Report	High	0.94	engineering	engineering	no	no (-)	PASS
S2	2	Feature Request	Feature Request	Low	0.94	product	product	no	no (-)	PASS
S3	2	Billing Issue	Billing Issue	Medium	0.95	billing	human_review_escalation	yes	yes (billing_amount_over_threshold)	PASS
S4	2	Technical Question	Technical Question	Low	0.94	technical_support	technical_support	either	no (-)	PASS
S5	2	Incident/Outage	Incident/Outage	High	0.96	incident_response	human_review_escalation	yes	yes (incident_category, outage_language)	PASS
S1	3	Bug Report	Bug Report	High	0.94	engineering	engineering	no	no (-)	PASS
S2	3	Feature Request	Feature Request	Low	0.94	product	product	no	no (-)	PASS
S3	3	Billing Issue	Billing Issue	Medium	0.95	billing	human_review_escalation	yes	yes (billing_amount_over_threshold)	PASS
S4	3	Technical Question	Technical Question	Low	0.94	technical_support	technical_support	either	no (-)	PASS
S5	3	Incident/Outage	Incident/Outage	High	0.96	incident_response	human_review_escalation	yes	yes (incident_category, outage_language)	PASS
E1-multi-intent	1	Billing Issue/Bug Report	Billing Issue	Medium	0.92	-	human_review_escalation	yes	yes (multi_intent)	PASS
E2-vague	1	Bug Report/Incident/Outage	Bug Report	Medium	0.78	-	engineering	yes	no (-)	FAIL: escalation, rules
E3-prompt-injection	1	Technical Question	Technical Question	Low	0.92	-	technical_support	either	no (-)	PASS
E4-french	1	Incident/Outage	Incident/Outage	High	0.96	-	human_review_escalation	yes	yes (incident_category)	PASS
E5-angry-small-billing	1	Billing Issue	Billing Issue	High	0.94	billing	billing	no	no (-)	FAIL: priority
E6-feature-sounds-like-bug	1	Feature Request/Bug Report	Bug Report	Medium	0.94	-	engineering	either	no (-)	PASS
E7-empty	1	(none)	None	High	None	-	human_review_escalation	yes	yes (llm_failure)	PASS
E8-long	1	Bug Report	Bug Report	Medium	0.94	engineering	engineering	no	no (-)	PASS

Consistency across runs

Case	Runs	Confidence range	Category / priority / destination / escalation
S1	3	0.94-0.94	stable
S2	3	0.94-0.94	stable
S3	3	0.95-0.95	stable
S4	3	0.94-0.94	stable
S5	3	0.96-0.96	stable